



Chapter 12: Spark Standalone Cluster

Ex1: NLP - Tags

Requirement: Build a tags filter. Use the various NLP tools and a classifier, to predict tag for one question. In future questions could be auto-tagged by such a classifier or tags could be recommended to users prior to posting.

- Dataset: stack-overflow-data.csv. It contains Stack Overflow questions and associated tags.
- Link tham khảo: <http://benalexkeen.com/multiclass-text-classification-with-pyspark/>
(<http://benalexkeen.com/multiclass-text-classification-with-pyspark/>)

```
In [1]: # Link HDFS: http://172.24.40.251:50070/
```

```
In [2]: import findspark
findspark.init()
```

```
In [3]: import pyspark
```

```
In [4]: from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark import SparkConf
```

```
In [5]: SparkContext.setSystemProperty('spark.executor.memory', '6g')

sc = SparkContext(master='spark://172.25.53.2:7077', appName='Stack_Overflow')
```

```
In [7]: sc
```

Out[7]: SparkContext

[Spark UI \(http://PM503-GV:4041\)](http://PM503-GV:4041)

Version

v2.4.4

Master

spark://172.25.53.2:7077

AppName

Stack_Overflow

```
In [8]: spark = SparkSession(sc)
```

```
In [10]: file_name = "hdfs://172.24.40.251:19000/stack_overflow_data.csv"
```



```
In [11]: data = spark.read.csv(file_name, inferSchema=True, header=True)
```

```
In [12]: data.show(5)
```

```
+-----+-----+
|          post|      tags|
+-----+-----+
|what is causing t...|      c#|
|have dynamic html...|  asp.net|
|how to convert a ...|objective-c|
|.net framework 4 ...|      .net|
|trying to calcula...|    python|
+-----+-----+
only showing top 5 rows
```

```
In [13]: data.groupby('tags').count().show(30)
```

```
+-----+-----+
|      tags|count|
+-----+-----+
|      iphone| 2000|
|      android| 2000|
|          c#| 2000|
|         null|20798|
|      asp.net| 2000|
|        html| 2000|
|       mysql| 2000|
|       jquery| 2000|
|  javascript| 2000|
|         css| 2000|
|         sql| 2000|
|        c++| 2000|
|          c| 2000|
|objective-c| 2000|
|        java| 2000|
|         php| 2000|
|        .net| 2000|
|         ios| 2000|
|       python| 2000|
|  angularjs| 2000|
|ruby-on-rails| 2000|
+-----+-----+
```

```
In [14]: tags_null_data = data.filter(data.tags.isNull())
```

```
In [15]: tags_null_data.count()
```

```
Out[15]: 20798
```

```
In [16]: data = data.filter(data.tags.isNotNull())
```



In [17]: `data.count()`

Out[17]: 40000

In [18]: `from pyspark.sql.functions import *`

Clean and Prepare the Data

**** Create a new length feature: ****

In [19]: `from pyspark.sql.functions import length`

In [20]: `data = data.withColumn('length',length(data['post']))`

In [21]: `data.show()`

```
+-----+-----+-----+
|          post|          tags|length|
+-----+-----+-----+
|what is causing t...|          c#|    833|
|have dynamic html...|       asp.net|    804|
|how to convert a ...| objective-c|    755|
|.net framework 4 ...|         .net|    349|
|trying to calcula...|       python|   1290|
|how to give alias...|       asp.net|    309|
|window.open() ret...|    angularjs|    495|
|identifying serve...|        iphone|    424|
|unknown method ke...|ruby-on-rails|   2022|
|from the include ...|    angularjs|   1279|
|when we need inte...|          c#|    995|
|how to install .i...|         ios|    344|
|dynamic textbox t...|       asp.net|    389|
|rather than bubbl...|          c|   1338|
|site deployed in ...|       asp.net|    349|
|connection in .ne...|         .net|    228|
|how to subtract 1...| objective-c|     62|
|ror console show ...|ruby-on-rails|   2594|
|distance between ...|        iphone|    336|
|sql query - how t...|         sql|   1037|
+-----+-----+-----+
only showing top 20 rows
```



```
In [22]: # Pretty Clear Difference
data.groupby('tags').mean().show()
```

```
+-----+-----+
|      tags|avg(length)|
+-----+-----+
|      iphone|      709.621|
|      android|     1713.4345|
|          c#|     1145.3065|
|    asp.net|      999.95|
|      html|     891.3105|
|      mysql|     1038.561|
|      jquery|     1081.507|
| javascript|      964.396|
|        css|      954.809|
|        sql|      870.912|
|       c++|     1295.955|
|         c|     1121.1115|
|objective-c|      972.8925|
|       java|     1357.308|
|       php|     1123.4205|
|       .net|      731.0075|
|        ios|      970.7565|
|      python|     1018.6695|
|  angularjs|     1294.7545|
|ruby-on-rails|     1244.2055|
+-----+-----+
```

Feature Transformations

```
In [23]: from bs4 import BeautifulSoup

from pyspark import keyword_only
from pyspark.ml import Transformer
from pyspark.ml.param.shared import HasInputCol, HasOutputCol
from pyspark.sql.functions import udf
from pyspark.sql.types import StringType
```



In [24]: `class BsTextExtractor(Transformer, HasInputCol, HasOutputCol):`

```
@keyword_only
def __init__(self, inputCol=None, outputCol=None):
    super(BsTextExtractor, self).__init__()
    kwargs = self._input_kwargs
    self.setParams(**kwargs)

@keyword_only
def setParams(self, inputCol=None, outputCol=None):
    kwargs = self._input_kwargs
    return self._set(**kwargs)

def _transform(self, dataset):

    def f(s):
        cleaned_post = BeautifulSoup(s).text
        return cleaned_post

    t = StringType()
    out_col = self.getOutputCol()
    in_col = dataset[self.getInputCol()]
    return dataset.withColumn(out_col, udf(f, t)(in_col))
```

In [25]: `from pyspark.ml.feature import Tokenizer, StopWordsRemover, CountVectorizer, IDF, StringIndexer`
`text_extractor = BsTextExtractor(inputCol="post", outputCol="cleaned_post")`
`tokenizer = Tokenizer(inputCol="cleaned_post", outputCol="token_text")`
`stopremove = StopWordsRemover(inputCol='token_text', outputCol='stop_tokens')`
`count_vec = CountVectorizer(inputCol='stop_tokens', outputCol='c_vec')`
`idf = IDF(inputCol="c_vec", outputCol="tf_idf")`
`class_to_num = StringIndexer(inputCol='tags', outputCol='label')`

In [26]: `from pyspark.ml.feature import VectorAssembler`
`from pyspark.ml.linalg import Vector`

In [27]: `clean_up = VectorAssembler(inputCols=['tf_idf', 'length'], outputCol='features')`

The Model

We'll use Naive Bayes, but feel free to play around with this choice!

In [28]: `from pyspark.ml.classification import NaiveBayes`

In [29]: `# Use defaults`
`nb = NaiveBayes()`

Pipeline

In [30]: `from pyspark.ml import Pipeline`

```
In [31]: data_prep_pipe = Pipeline(stages=[class_to_num,text_extractor,tokenizer,stopremov
```

```
In [32]: cleaner = data_prep_pipe.fit(data)
```

```
In [33]: clean_data = cleaner.transform(data)
```

Training and Evaluation!

```
In [34]: clean_data = clean_data.select(['label','features'])
```

```
In [35]: clean_data.show()
```

```
+-----+-----+
|label|          features|
+-----+-----+
| 14.0|(262145,[0,1,2,3,...|
|  3.0|(262145,[0,12,31,...|
| 16.0|(262145,[0,1,2,3,...|
|  8.0|(262145,[0,18,21,...|
|  9.0|(262145,[0,1,4,8,...|
|  3.0|(262145,[0,12,21,...|
| 11.0|(262145,[0,1,3,6,...|
| 18.0|(262145,[0,44,61,...|
|  6.0|(262145,[0,1,14,2...|
| 11.0|(262145,[0,1,3,4,...|
| 14.0|(262145,[0,2,3,6,...|
|  1.0|(262145,[0,18,27,...|
|  3.0|(262145,[0,7,12,1...|
| 13.0|(262145,[0,1,2,3,...|
|  3.0|(262145,[0,11,27,...|
|  8.0|(262145,[0,187,23...|
| 16.0|(262145,[0,10,15,...|
|  6.0|(262145,[0,1,3,12...|
| 18.0|(262145,[0,30,39,...|
|  0.0|(262145,[0,12,15,...|
+-----+-----+
```

only showing top 20 rows

```
In [36]: (training,testing) = clean_data.randomSplit([0.7,0.3], seed=142)
```

```
In [41]: predictor = nb.fit(training)
```

```
In [42]: test_results = predictor.transform(testing)
```





In [43]: test_results.show()

```

+-----+-----+-----+-----+-----+
-+
|label|          features|      rawPrediction|      probability|prediction
|
+-----+-----+-----+-----+-----+
-+
| 0.0|(262145,[0,1,4,11...|[-21169.668465935...|[1.21437232270286...|12.
0|
| 0.0|(262145,[0,1,4,14...|[-1167.4513441585...|[1.0,7.6094842671...|0.
0|
| 0.0|(262145,[0,1,5,14...|[-11564.690744964...|[1.0,0.0,0.0,0.0,...|0.
0|
| 0.0|(262145,[0,1,5,14...|[-10439.407817798...|[1.93390066856389...|12.
0|
| 0.0|(262145,[0,1,7,9,...|[-7336.5344958838...|[1.0,1.5922509763...|0.
0|
| 0.0|(262145,[0,1,7,11...|[-8486.4019708681...|[1.0,2.0825090588...|0.
0|
| 0.0|(262145,[0,1,9,10...|[-3738.5573599120...|[7.30286711565549...|12.
0|
| 0.0|(262145,[0,1,9,12...|[-6617.1798543878...|[0.99995740820888...|0.
0|
| 0.0|(262145,[0,1,9,12...|[-1693.4933675940...|[2.32251400846856...|12.
0|
| 0.0|(262145,[0,1,10,1...|[-3989.2579384985...|[1.0,3.3928388326...|0.
0|
| 0.0|(262145,[0,1,10,1...|[-2373.9652840133...|[1.0,2.7652251546...|0.
0|
| 0.0|(262145,[0,1,11,1...|[-8818.6290325138...|[1.0,1.4580653626...|0.
0|
| 0.0|(262145,[0,1,11,1...|[-4285.8899650084...|[1.0,9.2886878940...|0.
0|
| 0.0|(262145,[0,1,11,1...|[-2804.2071232226...|[1.0,3.5487447155...|0.
0|
| 0.0|(262145,[0,1,11,1...|[-2115.4985343601...|[6.48717059350829...|4.
0|
| 0.0|(262145,[0,1,12,1...|[-2828.9699626115...|[0.13219295903883...|12.
0|
| 0.0|(262145,[0,1,12,1...|[-5997.8626004770...|[1.0,1.4086114099...|0.
0|
| 0.0|(262145,[0,1,12,1...|[-3875.2807055205...|[1.0,2.0818433425...|0.
0|
| 0.0|(262145,[0,1,12,1...|[-5370.3720000707...|[1.0,2.0080477642...|0.
0|
| 0.0|(262145,[0,1,12,1...|[-7260.3672598605...|[1.0,1.6809702518...|0.
0|
+-----+-----+-----+-----+-----+
-+
only showing top 20 rows

```



```
In [44]: # Create a confusion matrix
test_results.groupBy('label', 'prediction').count().show()
```

```
+-----+-----+-----+
|label|prediction|count|
+-----+-----+-----+
| 8.0|      3.0|   53|
| 16.0|     8.0|   20|
| 19.0|     5.0|    1|
| 10.0|     1.0|    3|
| 12.0|     5.0|    2|
|  0.0|    12.0|  169|
|  0.0|     8.0|   17|
|  1.0|    19.0|    9|
|  1.0|    12.0|    1|
|  7.0|     3.0|   11|
| 15.0|    11.0|   12|
| 19.0|    12.0|    1|
| 17.0|     7.0|   12|
| 17.0|    19.0|   15|
| 11.0|    17.0|    7|
|  8.0|     6.0|    1|
| 17.0|     9.0|    6|
|  4.0|     6.0|    1|
|  3.0|     9.0|    3|
|  3.0|     5.0|    1|
+-----+-----+-----+
only showing top 20 rows
```

```
In [45]: from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
In [46]: acc_eval = MulticlassClassificationEvaluator()
acc = acc_eval.evaluate(test_results)
print("Accuracy of model at predicting: {}".format(acc))
```

Accuracy of model at predicting: 0.7148108175153408

```
In [48]: # save hdfs
nb.save("hdfs://172.24.40.251:19000/NB_TagFilters_model")
```

- Not very good result! (~72%)
- Solution: Try switching out the classification models! Or even try to come up with other engineered features!...

Use LogisticRegression/Random Forest

Logistic Regression

```
In [49]: from pyspark.ml.classification import RandomForestClassifier, LogisticRegression
```




```
In [50]: lg = LogisticRegression(maxIter=20, regParam=0.3, elasticNetParam=0)
```

```
In [51]: predictor_1 = lg.fit(training)
```

```
In [52]: test_results_1 = predictor_1.transform(testing)
```

```
In [53]: # Create a confusion matrix
test_results_1.groupBy('label', 'prediction').count().show()
```

```
+-----+-----+-----+
|label|prediction|count|
+-----+-----+-----+
| 8.0|      3.0|   47|
| 16.0|     8.0|   10|
| 19.0|     5.0|    1|
| 10.0|     1.0|    1|
|  2.0|     0.0|    2|
| 15.0|    16.0|    1|
| 12.0|     5.0|    3|
|  0.0|    12.0|   94|
|  0.0|     8.0|    8|
|  1.0|    19.0|   12|
|  1.0|    12.0|    3|
|  7.0|     3.0|    6|
| 15.0|    11.0|    5|
| 19.0|    12.0|    5|
| 11.0|    17.0|   15|
| 17.0|     7.0|   13|
| 17.0|    19.0|    2|
| 17.0|     9.0|   11|
|  8.0|     6.0|    5|
|  6.0|     1.0|    2|
+-----+-----+-----+
only showing top 20 rows
```

```
In [54]: acc_eval = MulticlassClassificationEvaluator()
acc_1 = acc_eval.evaluate(test_results_1)
print("Accuracy of model at predicting: {}".format(acc_1))
```

Accuracy of model at predicting: 0.7182544021412621

```
In [55]: ## It's not better result!!!
```

```
In [57]: # save hdfs
lg.save("hdfs://172.24.40.251:19000/LG_TagFilters_model")
```

Random forest



```
In [58]: rf = RandomForestClassifier(labelCol="label", \
                                   featuresCol="features", \
                                   numTrees = 500, \
                                   maxDepth = 5, \
                                   maxBins = 64)
```

```
In [59]: predictor_2 = rf.fit(training)
```

```
In [60]: test_results_2 = predictor_2.transform(testing)
```

```
In [61]: # Create a confusion matrix
test_results_2.groupBy('label', 'prediction').count().show()
```

```
+-----+-----+-----+
|label|prediction|count|
+-----+-----+-----+
| 8.0|      3.0|   33|
| 16.0|     8.0|   31|
| 19.0|     5.0|    1|
|  0.0|    12.0|  268|
|  1.0|    19.0|    5|
|  1.0|    12.0|   19|
|  0.0|     8.0|    3|
| 19.0|    12.0|   28|
|  7.0|     3.0|    1|
| 15.0|    11.0|    1|
| 11.0|    17.0|    4|
| 17.0|     7.0|   16|
|  8.0|     6.0|    4|
| 17.0|     9.0|    5|
|  6.0|     1.0|    1|
|  3.0|     9.0|    1|
|  4.0|     6.0|    1|
|  6.0|     8.0|    9|
| 15.0|    19.0|    1|
| 14.0|     7.0|    3|
+-----+-----+-----+
```

only showing top 20 rows



```
In [62]: test_results_2.groupBy('prediction').count().show()
```

```
+-----+-----+
|prediction|count|
+-----+-----+
|      8.0| 1046|
|      0.0|   375|
|      7.0|   582|
|     18.0|   426|
|      1.0|   524|
|      4.0|   762|
|     11.0|   580|
|     14.0|   251|
|      3.0|   491|
|     19.0|   396|
|      2.0|   647|
|     17.0|   402|
|     10.0|   624|
|     13.0|   720|
|      6.0|   578|
|     15.0|   605|
|      5.0|   520|
|      9.0|   592|
|     16.0|   656|
|     12.0|  1113|
+-----+-----+
```

```
In [63]: acc_eval = MulticlassClassificationEvaluator()
acc_2 = acc_eval.evaluate(test_results_2)
print("Accuracy of model at predicting: {}".format(acc_2))
```

Accuracy of model at predicting: 0.7231428382621157

```
In [64]: ## It has higher accuracy but is not a better result!!!
```

```
In [66]: # save hdfs
rf.save("hdfs://172.24.40.251:19000/RF_TagFilters_model")
```

```
In [67]: sc.stop()
```